

diamond-price-prediction

November 15, 2023

1 Problem Statements: Diamond Price Prediction

1.1 Description:

- Diamonds are the hardest naturally occurring substance known to man. They are made of pure carbon, arranged in a crystal structure called diamond cubic.
- Diamonds are found in volcanic rock, called kimberlite, which is formed when the Earth's mantle melts and rises to the surface.
- Diamonds can be colorless, or they can have a variety of colors, including yellow, brown, pink, blue, green, and red. The color of a diamond is determined by the impurities that are present in the carbon atoms.
- Diamonds are used in jewelry and in industrial applications. In jewelry, diamonds are prized for their beauty, durability, and rarity. In industry, diamonds are used for cutting, drilling, and polishing.
- The four Cs of diamonds are color, cut, clarity, and carat. These are the factors that determine the value of a diamond.

2 1.0. Importing Libraries

```
[3]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.impute import SimpleImputer ## Handling Missing Values
from sklearn.preprocessing import StandardScaler # Handling Feature Scaling
from sklearn.preprocessing import OrdinalEncoder # Ordinal Encoding
## pipelines
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
```

3 2.0. The Datasets

3.1 2.1. Reading Datasets

```
[4]: ## Data Ingestions step
df=pd.read_csv('data/gemstone.csv')
df.head()
```

```
[4]:   id  carat      cut color clarity depth  table     x     y     z  price
0   0    1.52  Premium     F    VS2   62.2   58.0  7.27  7.33  4.55  13619
1   1    2.03 Very Good     J    SI2   62.0   58.0  8.06  8.12  5.05  13387
2   2    0.70   Ideal     G    VS1   61.2   57.0  5.69  5.73  3.50   2772
3   3    0.32   Ideal     G    VS1   61.6   56.0  4.38  4.41  2.71    666
4   4    1.70  Premium     G    VS2   62.6   59.0  7.65  7.61  4.77  14453
```

3.2 2.2. Datasets Information:

3.2.1 Introduction About the Data :

The dataset The goal is to predict price of given diamond (Regression Analysis).

There are 10 independent variables (including id):

- **id** : unique identifier of each diamond
- **carat** : Carat (ct.) refers to the unique unit of weight measurement used exclusively to weigh gemstones and diamonds.
- **cut** : Quality of Diamond Cut
- **color** : Color of Diamond
- **clarity** : Diamond clarity is a measure of the purity and rarity of the stone, graded by the visibility of these characteristics under 10-power magnification.
- **depth** : The depth of diamond is its height (in millimeters) measured from the culet (bottom tip) to the table (flat, top surface)
- **table** : A diamond's table is the facet which can be seen when the stone is viewed face up.
- **x** : Diamond X dimension
- **y** : Diamond Y dimension
- **z** : Diamond Z dimension

Target variable: * **price**: Price of the given Diamond.

Dataset Source Link : <https://www.kaggle.com/competitions/playground-series-s3e8/data?select=train.csv>

4 3.0. Data Exploration

```
[5]: df.isnull().sum()
```

```
[5]: id          0
     carat       0
     cut        0
```

```

color      0
clarity    0
depth      0
table      0
x          0
y          0
z          0
price      0
dtype: int64

```

```
[6]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 193573 entries, 0 to 193572
Data columns (total 11 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   id          193573 non-null  int64
 1   carat       193573 non-null  float64
 2   cut         193573 non-null  object
 3   color       193573 non-null  object
 4   clarity     193573 non-null  object
 5   depth       193573 non-null  float64
 6   table       193573 non-null  float64
 7   x           193573 non-null  float64
 8   y           193573 non-null  float64
 9   z           193573 non-null  float64
10  price       193573 non-null  int64
dtypes: float64(6), int64(2), object(3)
memory usage: 16.2+ MB

```

```
[7]: df.head()
```

```

[7]:   id  carat      cut color clarity depth table     x     y     z  price
0   0   1.52   Premium     F     VS2   62.2   58.0  7.27  7.33  4.55  13619
1   1   2.03  Very Good     J     SI2   62.0   58.0  8.06  8.12  5.05  13387
2   2   0.70    Ideal     G     VS1   61.2   57.0  5.69  5.73  3.50   2772
3   3   0.32    Ideal     G     VS1   61.6   56.0  4.38  4.41  2.71    666
4   4   1.70   Premium     G     VS2   62.6   59.0  7.65  7.61  4.77  14453

```

- Lets drop the id column

```
[8]: df=df.drop(labels=['id'],axis=1)
df.head()
```

```

[8]:   carat      cut color clarity depth table     x     y     z  price
0   1.52   Premium     F     VS2   62.2   58.0  7.27  7.33  4.55  13619

```

1	2.03	Very Good	J	SI2	62.0	58.0	8.06	8.12	5.05	13387
2	0.70	Ideal	G	VS1	61.2	57.0	5.69	5.73	3.50	2772
3	0.32	Ideal	G	VS1	61.6	56.0	4.38	4.41	2.71	666
4	1.70	Premium	G	VS2	62.6	59.0	7.65	7.61	4.77	14453

- check for duplicated records

```
[9]: df.duplicated().sum()
```

```
[9]: 0
```

segregate numerical and categorical columns

```
[10]: numerical_columns=df.columns[df.dtypes!='object']
categorical_columns=df.columns[df.dtypes=='object']
print("Numerical columns:",numerical_columns)
print('Categorical Columns:',categorical_columns)
```

```
Numerical columns: Index(['carat', 'depth', 'table', 'x', 'y', 'z', 'price'],
dtype='object')
```

```
Categorical Columns: Index(['cut', 'color', 'clarity'], dtype='object')
```

```
[11]: df[categorical_columns].describe()
```

```
[11]:
```

	cut	color	clarity
count	193573	193573	193573
unique	5	7	8
top	Ideal	G	SI1
freq	92454	44391	53272

```
[12]: df['cut'].value_counts()
```

```
[12]: Ideal          92454
Premium          49910
Very Good       37566
Good            11622
Fair             2021
Name: cut, dtype: int64
```

```
[13]: df['color'].value_counts()
```

```
[13]: G      44391
E      35869
F      34258
H      30799
D      24286
I      17514
J       6456
```

Name: color, dtype: int64


```
[14]: df['clarity'].value_counts()
```


```
[14]: SI1      53272
      VS2      48027
      VS1      30669
      SI2      30484
      VVS2     15762
      VVS1     10628
      IF        4219
      I1         512
      Name: clarity, dtype: int64
```


5 4.0. Data Visualization


```
[15]: plt.style.use('dark_background')
```


```
[16]: plt.figure(figsize=(8,6))
      x=0
      for i in numerical_columns:
          sns.histplot(data=df,x=i,kde=True)
          print('\n')
          plt.show()
```


















```
[17]: ## correlation
sns.heatmap(df[numerical_columns].corr(),annot=True)
```

```
[17]: <AxesSubplot: >
```



6 5.0. Feature Selection

```
[18]: df.head()
```

```
[18]:   carat      cut color clarity  depth  table     x     y     z  price
0   1.52   Premium    F     VS2   62.2   58.0  7.27  7.33  4.55  13619
1   2.03  Very Good    J     SI2   62.0   58.0  8.06  8.12  5.05  13387
2   0.70    Ideal     G     VS1   61.2   57.0  5.69  5.73  3.50   2772
3   0.32    Ideal     G     VS1   61.6   56.0  4.38  4.41  2.71    666
4   1.70   Premium    G     VS2   62.6   59.0  7.65  7.61  4.77  14453
```

```
[19]: df['cut'].unique()
```

```
[19]: array(['Premium', 'Very Good', 'Ideal', 'Good', 'Fair'], dtype=object)
```

```
[20]: cut_map={"Fair":1,"Good":2,"Very Good":3,"Premium":4,"Ideal":5}
```

```
[21]: df['clarity'].unique()
```

```
[21]: array(['VS2', 'SI2', 'VS1', 'SI1', 'IF', 'VVS2', 'VVS1', 'I1'],
          dtype=object)
```

```
[22]: df['clarity'].unique()
```

```
[22]: array(['VS2', 'SI2', 'VS1', 'SI1', 'IF', 'VVS2', 'VVS1', 'I1'],  
        dtype=object)
```

```
[23]: df['clarity'].unique()
```

```
[23]: array(['VS2', 'SI2', 'VS1', 'SI1', 'IF', 'VVS2', 'VVS1', 'I1'],  
        dtype=object)
```

```
[24]: df['color'].unique()
```

```
[24]: array(['F', 'J', 'G', 'E', 'D', 'H', 'I'], dtype=object)
```

```
[25]: color_map = {"D":1 , "E":2 , "F":3 , "G":4 , "H":5 , "I":6, "J":7}
```

```
[27]: df.head()
```

```
[27]:   carat  cut color clarity depth  table     x     y     z  price  
0   1.52   4    F     VS2   62.2   58.0  7.27  7.33  4.55  13619  
1   2.03   3    J     SI2   62.0   58.0  8.06  8.12  5.05  13387  
2   0.70   5    G     VS1   61.2   57.0  5.69  5.73  3.50   2772  
3   0.32   5    G     VS1   61.6   56.0  4.38  4.41  2.71    666  
4   1.70   4    G     VS2   62.6   59.0  7.65  7.61  4.77  14453
```

7 5.0. Model Training

```
[32]: df2 = pd.read_csv('./data/gemstone.csv')  
df2.head()
```

```
[32]:   carat      cut color clarity depth  table     x     y     z  price  
0   1.52  Premium    F     VS2   62.2   58.0  7.27  7.33  4.55  13619  
1   2.03  Very Good    J     SI2   62.0   58.0  8.06  8.12  5.05  13387  
2   0.70    Ideal     G     VS1   61.2   57.0  5.69  5.73  3.50   2772  
3   0.32    Ideal     G     VS1   61.6   56.0  4.38  4.41  2.71    666  
4   1.70  Premium     G     VS2   62.6   59.0  7.65  7.61  4.77  14453
```

```
[33]: df2=df2.drop(labels=['id'],axis=1)  
      ## Independent and dependent features  
X = df.drop(labels=['price'],axis=1)  
Y = df[['price']]  
Y
```

```
[33]:   price  
0   13619  
1   13387
```

```

2          2772
3           666
4         14453
...
193568    1130
193569    2874
193570    3036
193571     681
193572    2258

```

```
[193573 rows x 1 columns]
```

7.0.1 Segregating and categorical variables

```
[34]: categorical_cols = X.select_dtypes(include='object').columns
      numerical_cols = X.select_dtypes(exclude='object').columns
```

7.0.2 Define the custom ranking for each ordinal variable

```
[35]: cut_categories = ['Fair', 'Good', 'Very Good', 'Premium', 'Ideal']
      color_categories = ['D', 'E', 'F', 'G', 'H', 'I', 'J']
      clarity_categories = ['I1', 'SI2', 'SI1', 'VS2', 'VS1', 'VVS2', 'VVS1', 'IF']
```

7.0.3 Numerical Pipeline

```
[36]: num_pipeline=Pipeline(
      steps=[
          ('imputer', SimpleImputer(strategy='median')),
          ('scaler', StandardScaler())
      ]
  )

  # Categorical Pipeline
  cat_pipeline=Pipeline(
      steps=[
          ('imputer', SimpleImputer(strategy='most_frequent')),
          ↪ ('ordinalencoder', OrdinalEncoder(categories=[cut_categories, color_categories, clarity_categories])),
          ('scaler', StandardScaler())
      ]
  )

  preprocessor=ColumnTransformer([
```

```
(('num_pipeline',num_pipeline,numerical_cols),
 ('cat_pipeline',cat_pipeline,categorical_cols)
])
```

7.0.4 Train test split

```
[37]: from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X,Y,test_size=0.
↪30,random_state=30)
```

```
[38]: X_train=pd.DataFrame(preprocessor.fit_transform(X_train),columns=preprocessor.
↪get_feature_names_out())
X_test=pd.DataFrame(preprocessor.transform(X_test),columns=preprocessor.
↪get_feature_names_out())
```

```
[39]: X_train.head()
```

```
[39]:   num_pipeline__carat  num_pipeline__depth  num_pipeline__table  \
0          -0.975439          -0.849607          -0.121531
1           0.235195           1.833637          -0.121531
2           0.494617           0.815855           0.399800
3          -1.018676           0.260701           0.921131
4          -0.953821          -0.664555          -0.642862

   num_pipeline__x  num_pipeline__y  num_pipeline__z  cat_pipeline__cut  \
0          -1.042757          -1.080970          -1.123150           0.874076
1           0.318447           0.279859           0.485354          -2.144558
2           0.570855           0.606458           0.673737          -0.132136
3          -1.214034          -1.244270          -1.195605          -0.132136
4          -1.069801          -1.044681          -1.094168           0.874076

   cat_pipeline__color  cat_pipeline__clarity
0           1.528722           1.352731
1          -0.935071          -0.646786
2           0.296826           0.686225
3           0.296826           0.019720
4           2.144670           1.352731
```

```
[40]: ## Model Training

from sklearn.linear_model import LinearRegression,Lasso,Ridge,ElasticNet
from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error
```

```
[41]: regression=LinearRegression()
regression.fit(X_train,y_train)
```



```
[41]: LinearRegression()
```

```
[42]: regression.coef_
```

```
[42]: array([[ 6433.66003594, -132.75843566, -70.42922179, -1720.30971463,  
          -499.29302619, -63.39317848,  72.44537247, -460.41604642,  
          650.76431652]])
```

```
[43]: regression.intercept_
```

```
[43]: array([3970.76628955])
```

```
[44]: def evaluate_model(true, predicted):  
    mae = mean_absolute_error(true, predicted)  
    mse = mean_squared_error(true, predicted)  
    rmse = np.sqrt(mean_squared_error(true, predicted))  
    r2_square = r2_score(true, predicted)  
    return mae, rmse, r2_square
```

7.0.5 Train multiple models

7.0.6 Model Ecaluation

```
[45]: models={  
    'LinearRegression':LinearRegression(),  
    'Lasso':Lasso(),  
    'Ridge':Ridge(),  
    'Elasticnet':ElasticNet()  
}  
trained_model_list=[]  
model_list=[]  
r2_list=[]  
  
for i in range(len(list(models))):  
    model=list(models.values())[i]  
    model.fit(X_train,y_train)  
  
    #Make Predictions  
    y_pred=model.predict(X_test)  
  
    mae, rmse, r2_square=evaluate_model(y_test,y_pred)  
  
    print(list(models.keys())[i])  
    model_list.append(list(models.keys())[i])  
  
    print('Model Training Performance')  
    print("RMSE:",rmse)
```

```

print("MAE:",mae)
print("R2 score",r2_square*100)

r2_list.append(r2_square)

print('='*35)
print('\n')

```

```

LinearRegression
Model Training Performance
RMSE: 1013.9047094344002
MAE: 674.025511579685
R2 score 93.68908248567512
=====

```

```

Lasso
Model Training Performance
RMSE: 1013.8784226767013
MAE: 675.0716923362162
R2 score 93.68940971841704
=====

```

```

Ridge
Model Training Performance
RMSE: 1013.9059272771628
MAE: 674.0555800798244
R2 score 93.6890673250594
=====

```

```

Elasticnet
Model Training Performance
RMSE: 1533.4162456064048
MAE: 1060.7368759154729
R2 score 85.56494831165182
=====

```

8 Reference:

[PWSKILLS](#)

9 Thank You!